

MLOps Lab Assingment - 1 — B22AI063

Name: Jyotin Goel

Roll Number: B22AI063

Email: b22ai063@iitj.ac.in

Links

- **Colab Notebook:**

[Open in Google Colab](#)

- **WANDB visualizations:**

[View Assignment on wandb](#)

- **GitHub Repository:**

[View Assignment on GitHub](#)

1. Objective

The objective of this lab was to:

- Train a Convolutional Neural Network (CNN) on the **CIFAR-10 dataset**
 - Implement a **custom PyTorch dataloader**
 - Compute **FLOPs** for the selected model
 - Train the model for **25–30 epochs**
 - Visualize:
 - **Gradient flow**
 - **Weight update flow**
 - Log all metrics and visualizations using **Weights & Biases (Wandb)**
 - Upload all code + report to GitHub (without pushing trained models)
-

2. Dataset

We used the **CIFAR-10 dataset**, which contains:

- 60,000 RGB images of size **32×32**
- 10 classes (airplane, car, bird, cat, etc.)
- Train/Test split:
 - 50,000 training images
 - 10,000 test images

3. Model Architecture

A CNN model (**ResNet-18**) was chosen due to its strong performance and efficiency on CIFAR-10.

Key components:

- Convolutional layers with residual blocks
 - Batch normalization
 - ReLU activations
 - Fully connected classifier head
-

4. Custom DataLoader Implementation

Instead of using the default torchvision dataloader directly, a **custom Dataset class** was implemented.

Features:

- Custom `__getitem__` and `__len__`
 - Data augmentation:
 - Random crop
 - Horizontal flip
 - Normalization for CIFAR-10 statistics
-

5. FLOPs Computation

FLOPs (Floating Point Operations) were computed using the `ptflops` library.

Example output:

- **Total FLOPs:** ~1.8 GFLOPs
- **Total Parameters:** ~11.2M

This provides an estimate of computational cost during inference/training.

6. Training Setup

Training was performed for **30 epochs** with the following configuration:

Hyperparameter	Value
Optimizer	Adam
Learning Rate	0.001
Batch Size	128
Epochs	30

Hyperparameter	Value
Loss	CrossEntropyLoss

7. Gradient Flow Visualization

Gradient flow plots were generated during training to inspect whether gradients vanish or explode.

Observations:

- Early convolution layers had smaller gradients
- Deeper layers showed stronger gradient updates
- No major gradient instability was observed

Gradient flow plots were logged to Wandb each epoch.

8. Weight Update Flow Visualization

Weight update flow tracks how much parameters change over epochs.

Observations:

- Initial epochs showed large weight updates
- Updates stabilized after ~15 epochs
- Indicates convergence of training

These plots were also logged to Wandb.

9. Results and Findings

Training Performance

- Training accuracy steadily increased over epochs
- Test accuracy improved and stabilized around final epochs

10. Wandb Visualizations Logged

The following were logged to Wandb:

- Training loss curve
- Validation accuracy curve
- Gradient flow plots
- Weight update flow plots
- FLOPs + parameter summary